

113-1 (Fall 2024) Semester

Reinforcement Learning

Lecture 4: Model-Free Prediction

Slides




Shao-Hua Sun (孫紹華)

Assistant Professor in Electrical Engineering,
National Taiwan University

Disclaimer

Experimental course

- Everything is subject to change
- Most materials are made from scratch for this course
 - There could be mistakes and flaws in slides, assignments, etc.
- We are altogether in an early iteration of gradient descent
 - TAs and I are the optimizer, you (and your feedback) are the data




[[source](#)]

Credit to David Silver

- This part of the course — Introduction to Reinforcement Learning — is 99% based on [David Silver's Reinforcement Learning lecture](#) at University College London with DeepMind
- You are highly encouraged to watch David's lectures

Google DeepMind

[Research](#) [Blog](#) [Impact](#) [Safety & Ethics](#) [About](#) [Careers](#) 



Introduction to Reinforcement Learning with David Silver

Schedule

Week	Date	Topic	Assignment	Project
1	9/5	<ul style="list-style-type: none">• Lecture 0: Course Introduction• Lecture 1: Introduction to Reinforcement Learning		
2	9/12	<ul style="list-style-type: none">• Lecture 2: Markov Decision Processes• Lecture 3: Planning by Dynamic Programming	#1 Release	
3	9/19	<ul style="list-style-type: none">• Lecture 4: Model-Free Prediction• Lecture 5: Model-Free Control		
4	9/26	<ul style="list-style-type: none">• Lecture 6: Value Function Approximation• Lecture 7: Policy Gradient Methods	#1 Due #2 Release	

- Lecture 4 Model-Free Prediction (9:30 AM-10:50 AM)
- Lecture 5 Model-Free Control (11 AM - 12:10 PM)

Schedule

Week	Date	Topic	Assignment	Project
1	9/5	<ul style="list-style-type: none">Lecture 0: Course IntroductionLecture 1: Introduction to Reinforcement Learning		
2	9/12	<ul style="list-style-type: none">Lecture 2: Markov Decision ProcessesLecture 3: Planning by Dynamic Programming	#1 Release	
3	9/19	<ul style="list-style-type: none">Lecture 4: Model-Free PredictionLecture 5: Model-Free Control		
4	9/26	<ul style="list-style-type: none">Lecture 6: Value Function ApproximationLecture 7: Policy Gradient Methods	#1 Due #2 Release	

- Assignment #1
 - Deadline: 9/26 9:30 AM (no late submission)
- Assignment #2 TA session (9/26 11:40 AM - 12:10 PM) by 楊可 Co Yong
 - Release: 9/26 12:10 PM on NTU COOL
 - Deadline: 10/17 9:30 AM (no late submission)

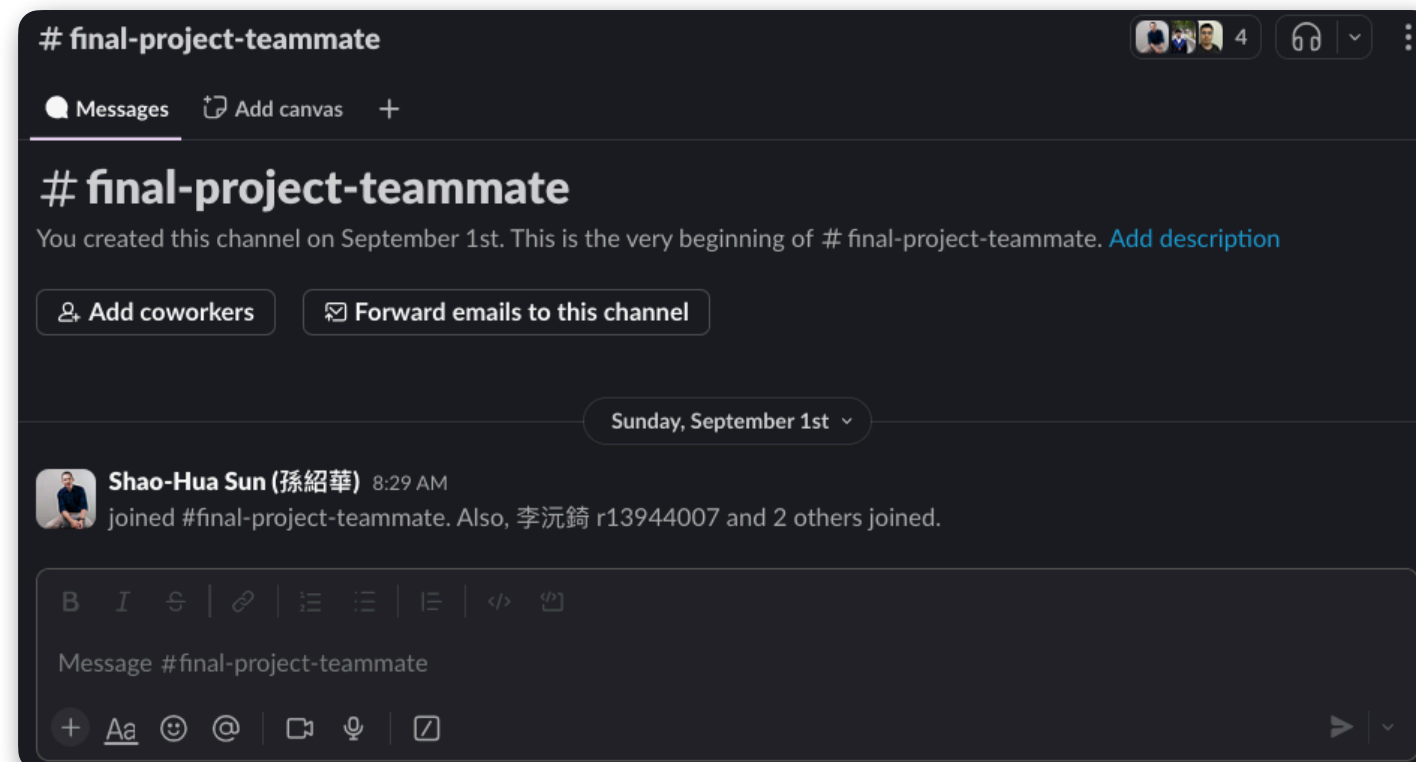
Final Project - Form Your Team

Form a team with 4 members

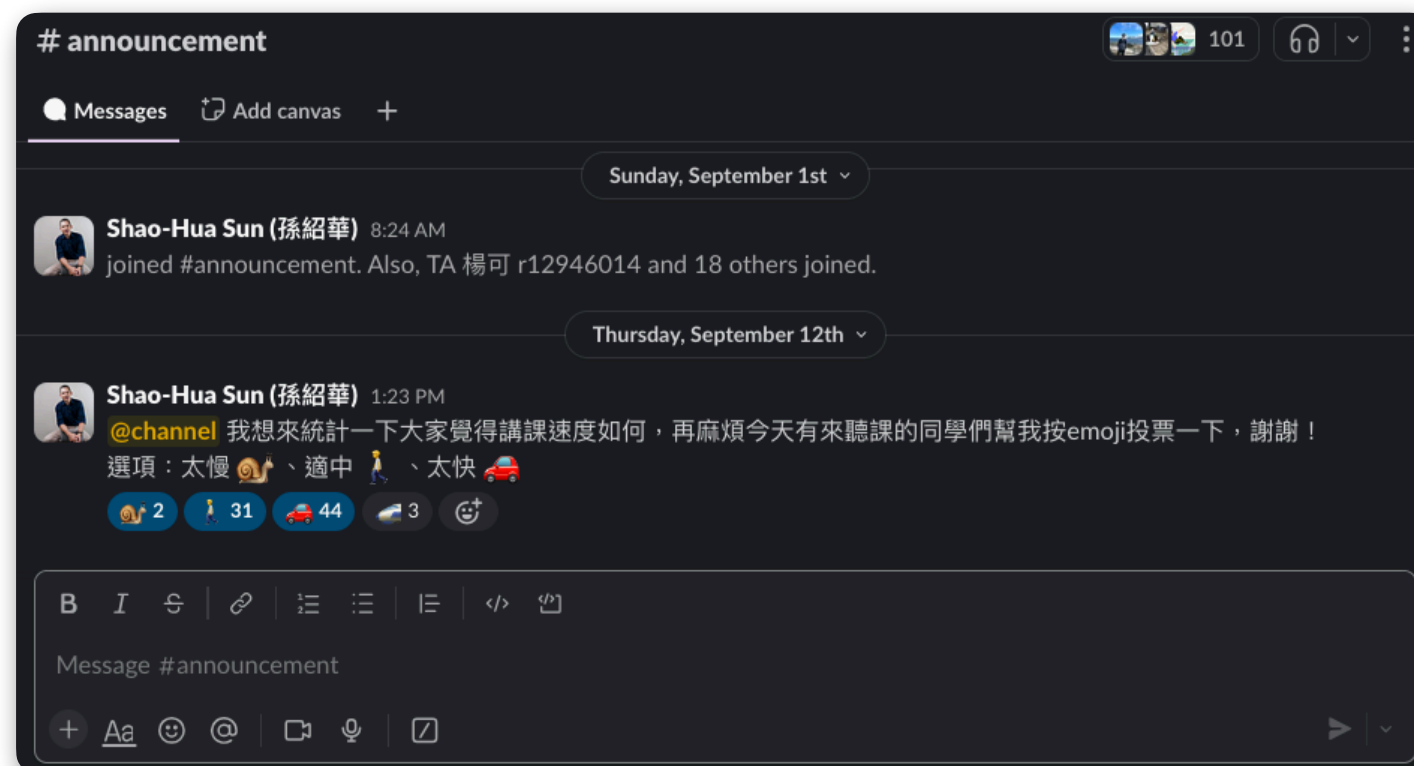
- Friends
- Lab mates
- Find them on slack `#final-project-teammate`

Based on

- Research interests
- Work habits



The Pace of Lectures



Week	Date	Topic	Assignment	Project
1	9/5	<ul style="list-style-type: none"> Lecture 0: Course Introduction Lecture 1: Introduction to Reinforcement Learning 		
2	9/12	<ul style="list-style-type: none"> Lecture 2: Markov Decision Processes Lecture 3: Planning by Dynamic Programming 	#1 Release	
3	9/19	<ul style="list-style-type: none"> Lecture 4: Model-Free Prediction Lecture 5: Model-Free Control 		
4	9/26	<ul style="list-style-type: none"> Lecture 6: Value Function Approximation Lecture 7: Policy Gradient Methods 	#1 Due #2 Release	
5	10/3	<ul style="list-style-type: none"> Lecture 8: Integrating Learning and Planning Lecture 9: Exploration and Exploitation 		
6	10/10	<ul style="list-style-type: none"> Lecture 10: Deep Q-Learning Lecture 11: Deep Policy Optimization 		
7	10/17	<ul style="list-style-type: none"> Lecture 11: Deep Policy Optimization Lecture 12: Deep Q-Learning + Policy Optimization 	#2 Due #3 Release	
8	10/24	<ul style="list-style-type: none"> Lecture 13: Imitation Learning Lecture 14: Skill-based RL Lecture 15: Offline RL 		

Week	Date	Topic	Assignment	Project
9	10/31	<ul style="list-style-type: none"> Lecture 16: Multi-task RL Lecture 17: Meta RL Lecture 18: Hierarchical RL 		Confirm team members and potential topics
10	11/7	<ul style="list-style-type: none"> Lecture 19: RL Exploration Lecture 20: Model-based RL Lecture 21: Programmatic RL Lecture 22: RL from Human Feedback 	#3 Due	
11	11/14	<ul style="list-style-type: none"> Final Project Proposal 		Meet with TA
12	11/21	<ul style="list-style-type: none"> Jiayuan Mao (MIT) Karl Pertsch (UC Berkeley & Stanford) 		Meet with the instructor
13	11/28	<ul style="list-style-type: none"> Youngwoon Lee (UC Berkeley) Guanzhi Wang (Caltech & Nvidia) 		Meet with TA
14	12/5	<ul style="list-style-type: none"> Risto Vuorio (University of Oxford) Kuang-Huei Lee (Google DeepMind) 		Meet with the instructor
15	12/12	<ul style="list-style-type: none"> Aleksei Petrenko (Apple) Ping-Chun Hsieh (NYCU) 		Meet with TA
16	12/19	<ul style="list-style-type: none"> Final Project Presentation 		Report deadline (12/22 11:59 PM)

Recap

Markov Process and its Variants

Category	Reward	Action	Problem
Markov Process (Markov Chain)	✗	✗	
Markov Reward Process (MRP)	✓	✗	Prediction
Markov Decision Process (MDP)	✓	✓	Prediction & Control

Bellman Equation - Summary

Bellman expectation equations

- State-value function v_π

$$v_\pi(s) = \sum_{a \in A} \pi(a | s) \left(R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_\pi(s') \right)$$

- Action-value function q_π

$$q_\pi(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \sum_{a' \in A} \pi(a' | s') q_\pi(s', a')$$

Bellman optimality equations

- Optimal state-value function v_*

$$v_*(s) = \max_a R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_*(s')$$

- Optimal action-value function q_*

$$q_*(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \max_{a'} q_*(s', a')$$

Iterative Policy Evaluation


Policy evaluation

- Problem: evaluate a given policy π
- Solution: iteratively apply Bellman expectation backup
 - v_1 (arbitrarily initialized) $\rightarrow v_2 \rightarrow v_3 \rightarrow \dots \rightarrow v_\pi$

Policy evaluation procedure

- At each iteration $k + 1$
- For all states $s \in S$
- Update $v_{k+1}(s)$ from $v_k(s')$ by

$$v_{k+1}(s) = \sum_{a \in A} \pi(a | s) \left(R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_k(s') \right) \quad \text{or} \quad v_{k+1} = R^\pi + \gamma P^\pi v_k$$



Convergence proof by the contraction mapping theorem: [reference](#)



Policy Iteration

Given policy π

- Evaluate the policy
- Improve the policy by acting greedily with respect to v_π , $\pi' = \text{greedy}(v_\pi)$

Policy iteration

- Policy evaluation
 - Estimate v_π
 - Iterative policy evaluation
- Policy improvement
 - Generate $\pi' \geq \pi$
 - Greedy policy improvement



Gridworld: the improved policy was optimal, $\pi' = \pi_*$ when $k = 3$

- In general, it needs more iterations of improvement / evaluation
 - This process of policy iteration always converges to π_*

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

	←	←	↙
↑	↖	↖	↓
↑	↗	↘	↓
↘	→	→	

Value Iteration


Value iteration

- Problem: find an optimal policy π
- Solution: iteratively apply Bellman optimality backup
 - v_1 (arbitrarily initialized) $\rightarrow v_2 \rightarrow v_3 \rightarrow \dots \rightarrow v_*$ (c.f., v_π in iterative policy evaluation)

Value iteration procedure

- At each iteration $k + 1$
- For all states $s \in S$
- Update $v_{k+1}(s)$ from $v_k(s')$ by

$$v_{k+1}(s) = \max_{a \in A} \left(R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_k(s') \right) \quad \text{or} \quad v_{k+1} = \max_{a \in A} R^a + \gamma P^a v_k$$



Convergence proof by the contraction mapping theorem: [reference](#)

Unlike policy iteration, there is no explicit policy

- Intermediate value functions may not correspond to any policy

Synchronous Dynamic Programming

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

Complexity

- Algorithms are based on state-value function $v_{\pi}(s)$ or $v_{*}(s)$
 - Complexity $O(mn^2)$ per iteration for m actions and n states
- Could also apply to action-value function $q_{\pi}(s, a)$ or $q_{*}(s, a)$
 - Complexity $O(m^2n^2)$ per iteration

Outline

- Introduction
- Monte-Carlo Learning
- Temporal-Difference Learning
- TD(λ)


Outline

- Introduction
- Monte-Carlo Learning
- Temporal-Difference Learning
- TD(λ)

Model-Free Reinforcement Learning

Last lecture (lecture 3)

- Problem: solve a **known** MDP
 - Prediction: policy evaluation
 - Control: policy iteration and value iteration
- Solution: planning by dynamic programming



This lecture (lecture 4)

- Problem: estimate the value function of an **unknown** MDP
 - Model-free prediction

Next lecture (lecture 5)

- Problem: optimize the value function and/or policy of an **unknown** MDP
 - Model-free control



Outline

- Introduction
- **Monte-Carlo Learning**
- Temporal-Difference Learning
- TD(λ)

Monte-Carlo Learning

Monte-Carlo (MC) reinforcement learning

- MC methods learn directly from episodes of experience
- MC is **model-free**: no knowledge of MDP transitions / rewards
- MC learns from **complete episodes**: no bootstrapping
- MC uses the simplest possible idea: value = mean return
- **Caveat**: can only apply MC to episodic MDPs
 - All episodes must terminate

Monte-Carlo policy evaluation

- Goal: learn v_π from **episodes** of experience under policy π ,

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

- The return is the total discounted reward: $G_t = R_{t+1} + \gamma R_{t+2} + \dots \gamma^{T-1} R_T$
- The value function is the expected return: $v_\pi(s) = \mathbb{E}[G_t | S_t = s]$
- Monte-Carlo policy evaluation uses *empirical mean return* instead of expected return

Monte-Carlo Policy Evaluation

First-visit Monte-Carlo policy evaluation

- To evaluate state s , the **first** time-step t that state s is visited in an episode
 - Increment counter $N(s) \leftarrow N(s) + 1$
 - Increment total return $Returns(s) \leftarrow Returns(s) + G_t$
 - Value is estimated by mean return $V(s) = Returns(s)/N(s)$
- By law of large numbers, $V(s) \rightarrow v_\pi(s)$ as $N(s) \rightarrow \infty$

Every-visit Monte-Carlo policy evaluation

- To evaluate state s , **every** time-step t that state s is visited in an episode
 - Increment counter $N(s) \leftarrow N(s) + 1$
 - Increment total return $Returns(s) \leftarrow Returns(s) + G_t$
 - Value is estimated by mean return $V(s) = Returns(s)/N(s)$
- Again, $V(s) \rightarrow v_\pi(s)$ as $N(s) \rightarrow \infty$

Example - Blackjack

Blackjack MDP

- States (200 of them)
 - Current sum (12-21), dealer's showing card (ace-10), do I have a “useable” ace? (yes/no)
- Actions
 - **stand/stick**: stop receiving cards (and terminate)
 - **hit/twist**: take another card (no replacement)
- Reward
 - **stand/stick**: stop receiving cards (and terminate)
 - +1 if sum of cards > sum of dealer cards
 - 0 if sum of cards = sum of dealer cards
 - -1 if sum of cards < sum of dealer cards
 - **hit/twist**: take another card (no replacement)
 - -1 if sum of cards > 21 (and terminate), 0 otherwise
- Transitions: automatically **hit** if sum of cards < 12

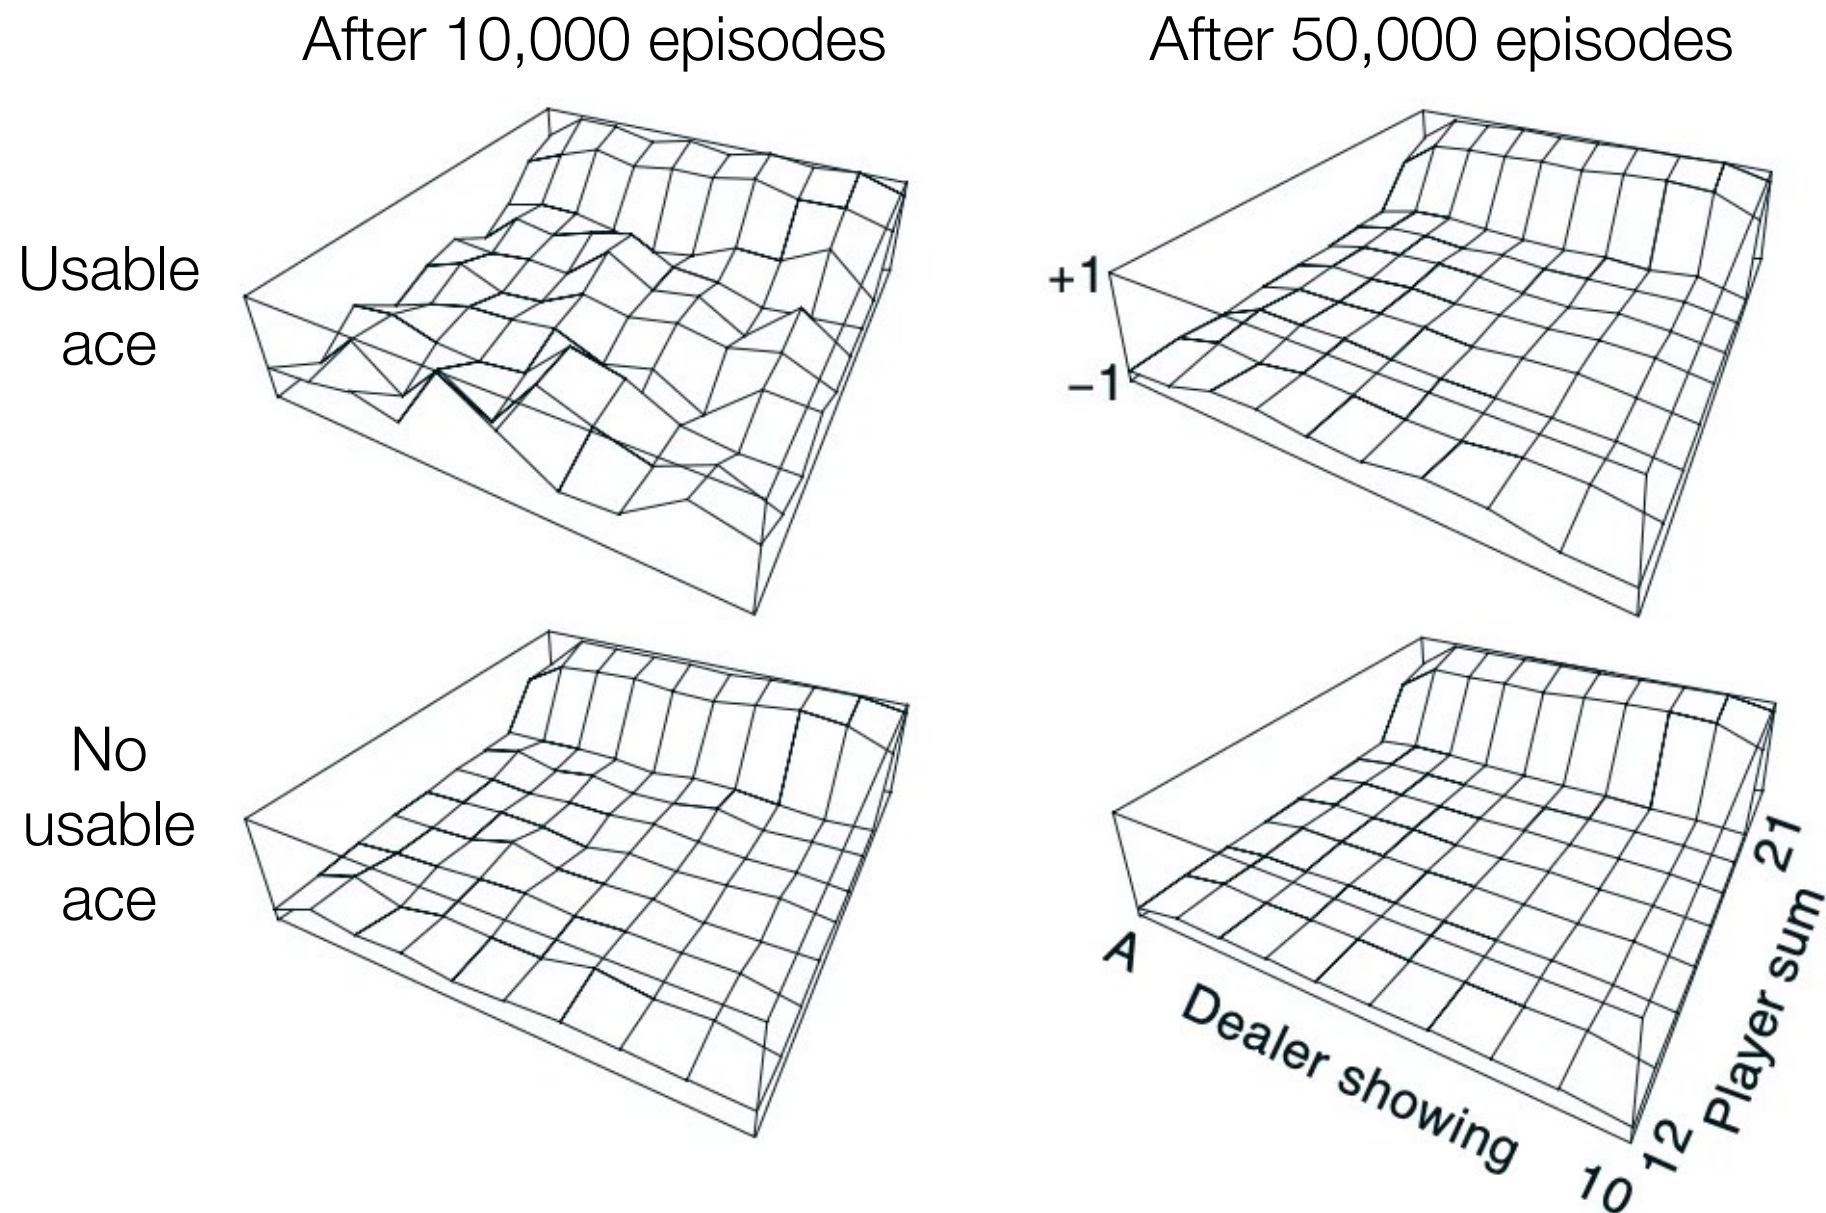


[\[source\]](#)

Blackjack Value Function with Monte-Carlo

Blackjack MDP

- Policy: **stand** if sum of cards ≥ 20 , otherwise **hit**



Incremental Monte-Carlo Updates

Incremental mean

- The mean μ_1, μ_2, \dots of a sequence x_1, x_2, \dots can be computed incrementally,

$$\mu_k = \frac{1}{k} \sum_{t=1}^k x_t = \frac{1}{k} \left(x_k + \sum_{t=1}^{k-1} x_t \right) = \frac{1}{k} [x_k + (k-1)\mu_{k-1}] = \mu_{k-1} + \frac{1}{k} (x_k - \mu_{k-1})$$

Incremental Monte-Carlo updates

- Update $V(s)$ incrementally after episodes $S_1, A_1, R_2, \dots, S_T$
- For each state S_t with return G_t
 - Increment counter $N(s) \leftarrow N(s) + 1$
 - Update the value $V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t))$
- In non-stationary problems, it can be useful to track a running mean, i.e., forget old episodes, $V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$, where α is the step size

Outline

- Introduction
- Monte-Carlo Learning
- Temporal-Difference Learning
- TD(λ)

Temporal-Difference Learning

Temporal-difference (TD) learning

- TD methods learn directly from episodes of experience
- TD is **model-free**: no knowledge of MDP transitions / rewards
- TD learns from **incomplete episodes**, by **bootstrapping**
- TD updates a guess towards a guess

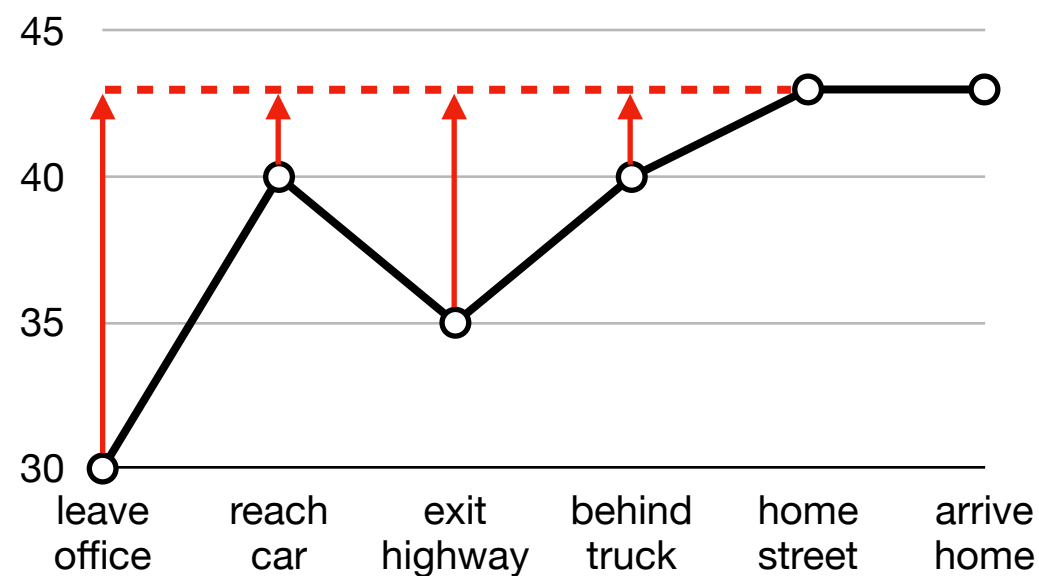
MC and TD: both learn v_π from **episodes** of experience under policy π

- Incremental every-visit Monte-Carlo: learn from a complete episode
 - Update value $V(S_t)$ toward *actual* return G_t by $V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$
- Simplest temporal-difference learning algorithm: TD(0)
 - Update value $V(S_t)$ toward *estimated* return $R_{t+1} + \gamma V(S_{t+1})$
$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$
 - $R_{t+1} + \gamma V(S_{t+1})$ is called the TD target
 - $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is called the TD error


Example - Driving Home

State	Elapsed Time	Predicted Time to Go	Predicted Total Time
Leave office	0	30	30
Reach car, raining	5	35	40
Exit highway	20	15	35
Behind truck	30	10	40
Home street	40	3	43
Arrive home	43	0	43

Changes recommended
by MC methods



Changes recommended
by TD methods



MC vs. TD - Final Outcome

MC backup

$$V(S_t) \leftarrow V(S_t) + \alpha(\textcolor{red}{G}_t - V(S_t))$$

TD backup

$$V(S_t) \leftarrow V(S_t) + \alpha(\textcolor{red}{R}_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

TD can learn **before** knowing the final outcome

- TD can learn online after every step
- MC must wait until end of episode before return is known

TD can learn **without** the final outcome

- TD can learn from incomplete sequences
- MC can only learn from complete sequences
- TD works in continuing (non-terminating) environments
- MC only works for episodic (terminating) environments

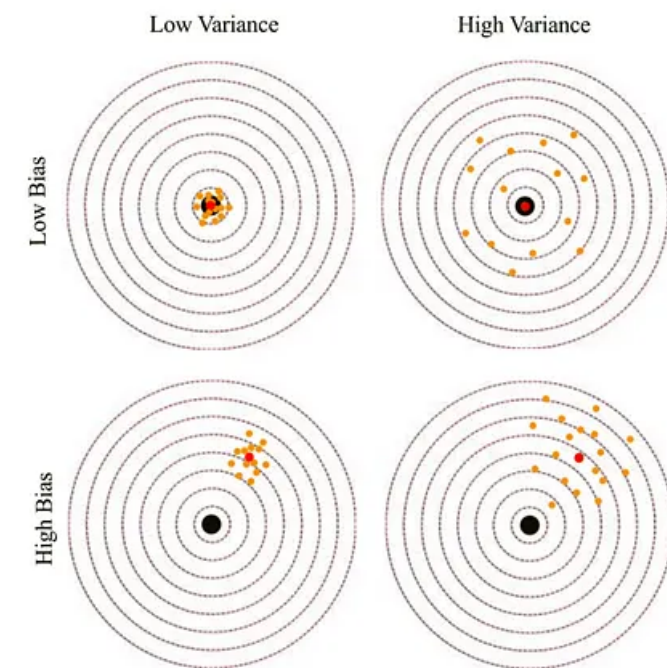
MC vs. TD - Bias/Variance Trade-Off

MC has **high variance, zero bias**

- Return $G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$ is **unbiased estimate** of $v_\pi(S_t)$
- Good convergence properties (even with function approximation)
- Not very sensitive to initial value
- Very simple to understand and use

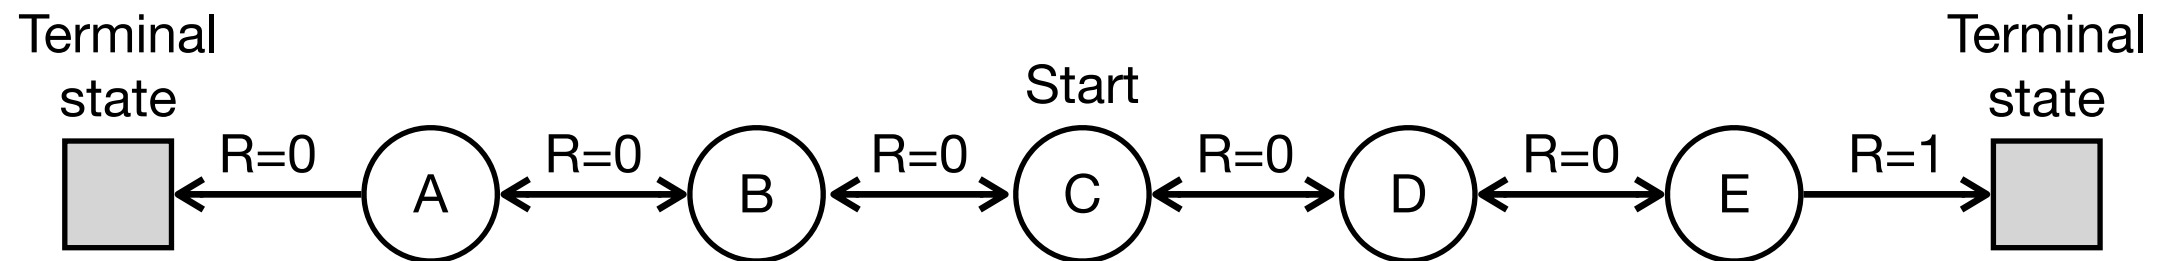
TD has **low variance, some bias**

- True TD target $R_{t+1} + \gamma v_\pi(S_{t+1})$ is **unbiased estimate** of $v_\pi(S_t)$
 - TD target $R_{t+1} + \gamma V(S_{t+1})$ is **biased estimate** of $v_\pi(S_t)$
- TD target is much **lower variance** than the return
 - Return depends on many random actions, transitions, rewards
 - TD target depends on one random action, transition, reward
- Usually more efficient than MC
- TD(0) converges to $v_\pi(s)$ (but not always with function approximation)
- More sensitive to initial value




[\[source\]](#)

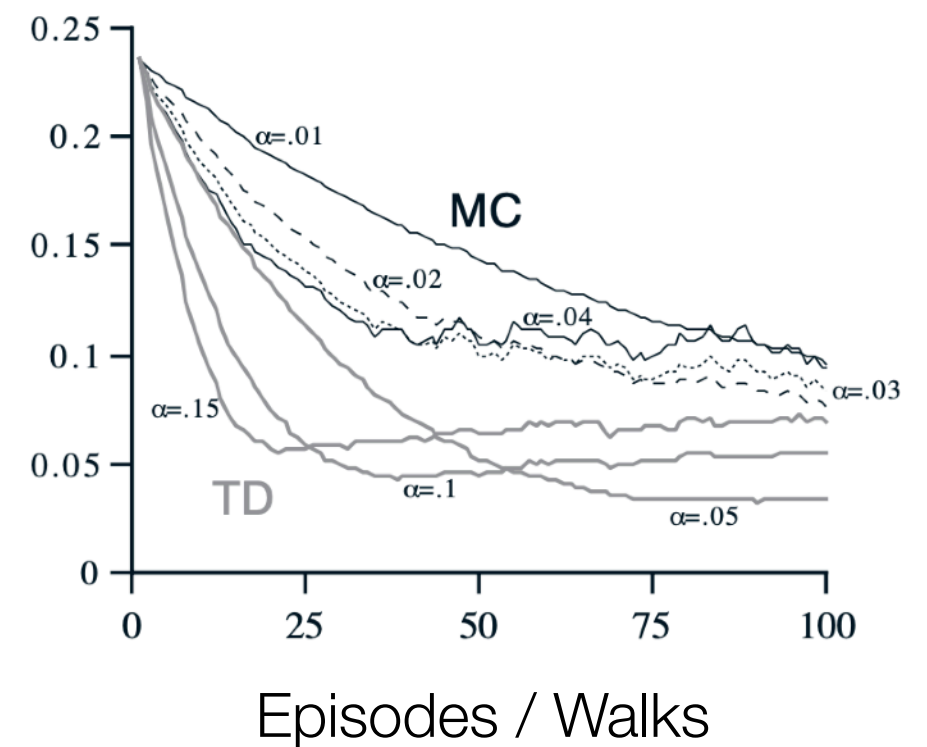
Example - Random Walk



Policy: agent follows a uniform random policy, $\pi(\rightarrow | \cdot) = \pi(\leftarrow | \cdot) = 0.5$



RMS error,
averaged
over states



Batch MC and TD

MC and TD converge: $V(s) \rightarrow v_\pi(s)$ as experience $\rightarrow \infty$

- But what about batch solution for finite experience?

Episode 1 $s_1^1, a_1^1, r_2^1, \dots, s_{T_1}^1$
 \vdots
 Episode K $s_1^K, a_1^K, r_2^K, \dots, s_{T_1}^K$

- Repeatedly sample episode $k \in [1, K]$ and apply MC or TD(0) to episode k


Example: two states A, B ; no discounting; 8 episodes of experience; what is $V(A), V(B)$?

- A, 0, B, 0
- B, 1
- B, 1
- B, 1
- B, 1
- B, 1
- B, 1
- B, 0



Value	TD	MC
$V(A)$	0.75	0
$V(B)$	0.75	0.75

Underlying MDP



MC vs. TD - Exploit Markov Property

MC converges to solution with *minimum mean-squared error*

- Best fit to the observed returns

$$\sum_{k=1}^K \sum_{t=1}^{T_k} (G_t^k - V(s_t^k))^2$$

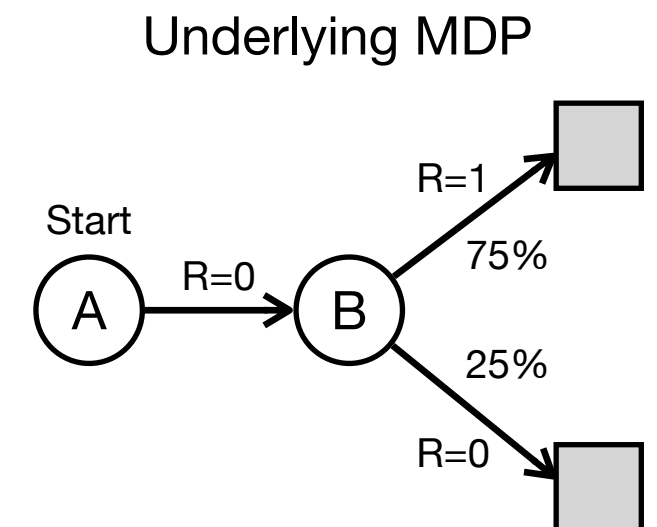
- In the AB example, $V(A) = 0$
- MC does not exploit Markov property: usually more effective in non-Markov environments

TD(0) converges to solution of *max likelihood Markov model*

- Solution to the MDP $\langle S, A, \hat{P}, \hat{R}, \gamma \rangle$ that best fits the data

$$\hat{P}_{s,s'}^a = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} \mathbf{1}(s_t^k, a_t^k, s_{t+1}^k = s, a, s')$$
$$\hat{R}_s^a = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} \mathbf{1}(s_t^k, a_t^k = s, a) r_t^k$$


- In the AB example, $V(A) = 0.75$
- TD exploits Markov property: usually more efficient in Markov environments



Comparison - MC, TD, and DP


MC backup

$$V(S_t) + \alpha(G_t - V(S_t))$$




TD backup

$$V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



DP backup

$$\mathbb{E}_{\pi}[R_{t+1} + \gamma V(S_{t+1})]$$



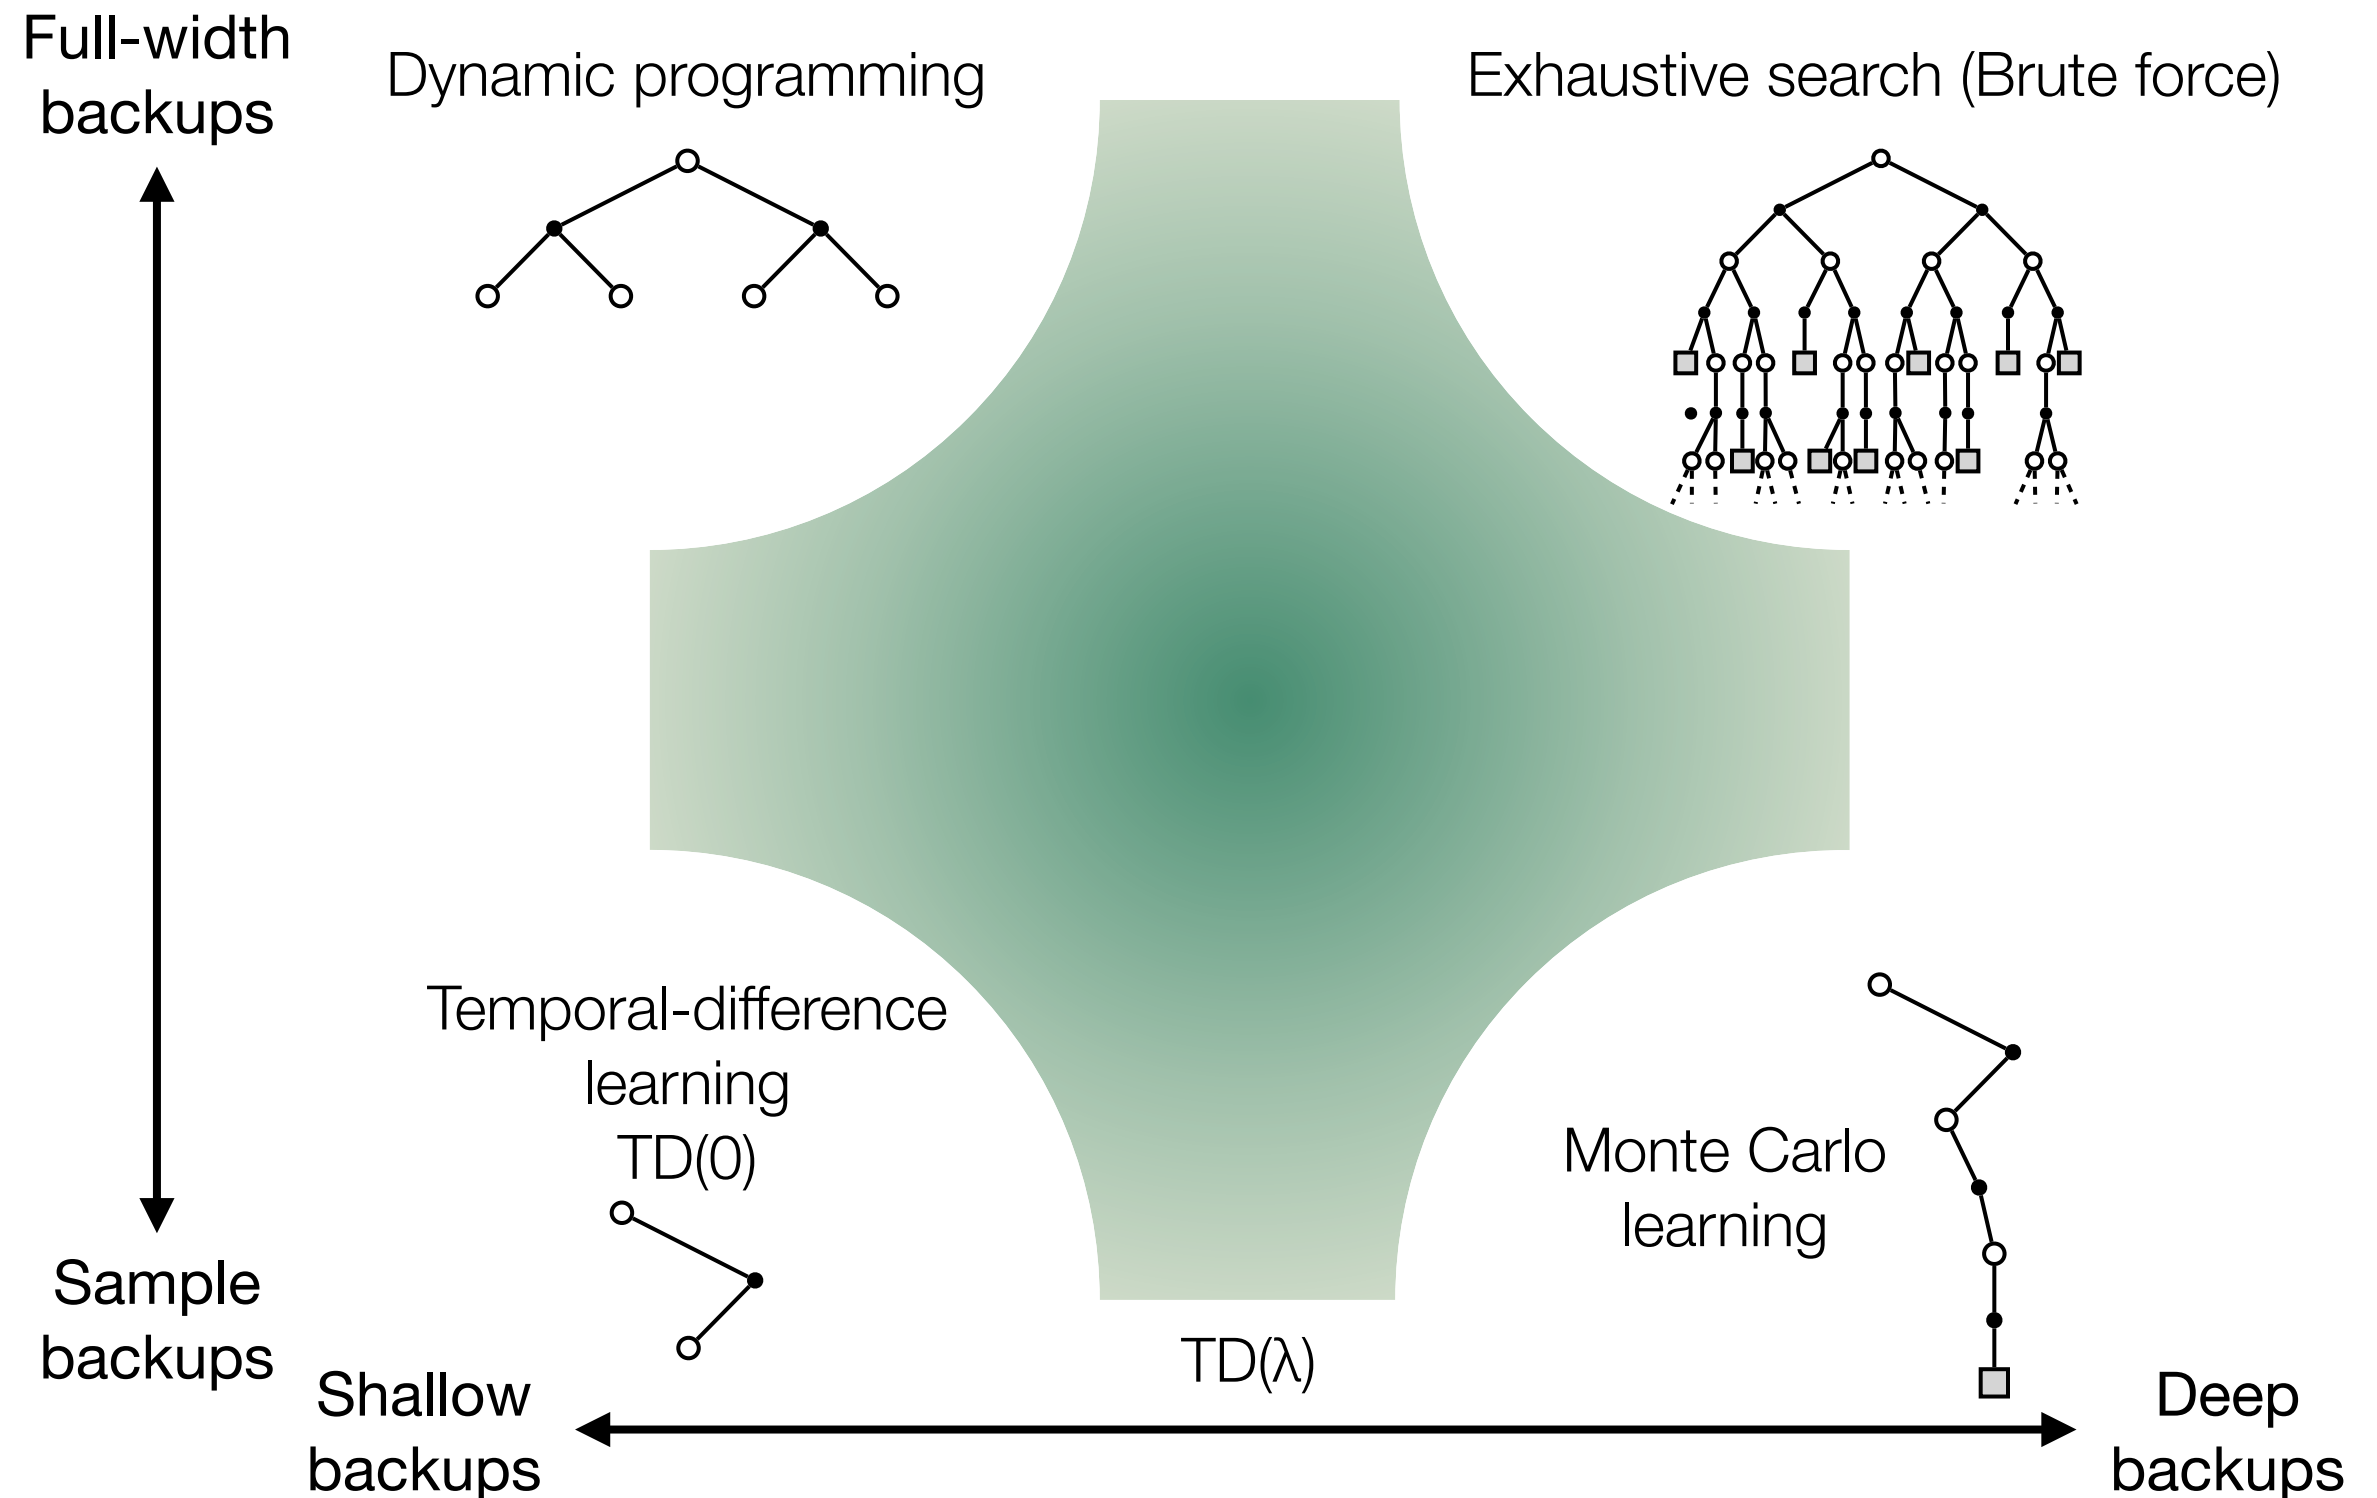
Bootstrapping: update involves an **estimate**

- MC does not bootstrap; TD and DP bootstrap

Sampling: update **samples** an expectation

- MC and TD sample; DP does not sample

Comparison - MC, TD, and DP




Outline

- Introduction
- Monte-Carlo Learning
- Temporal-Difference Learning
- TD(λ)

n-Step Prediction and Return

n-step Prediction: Let TD target
look n steps into the future

Consider the following n-step returns
for $n = 1, 2, \dots, \infty$

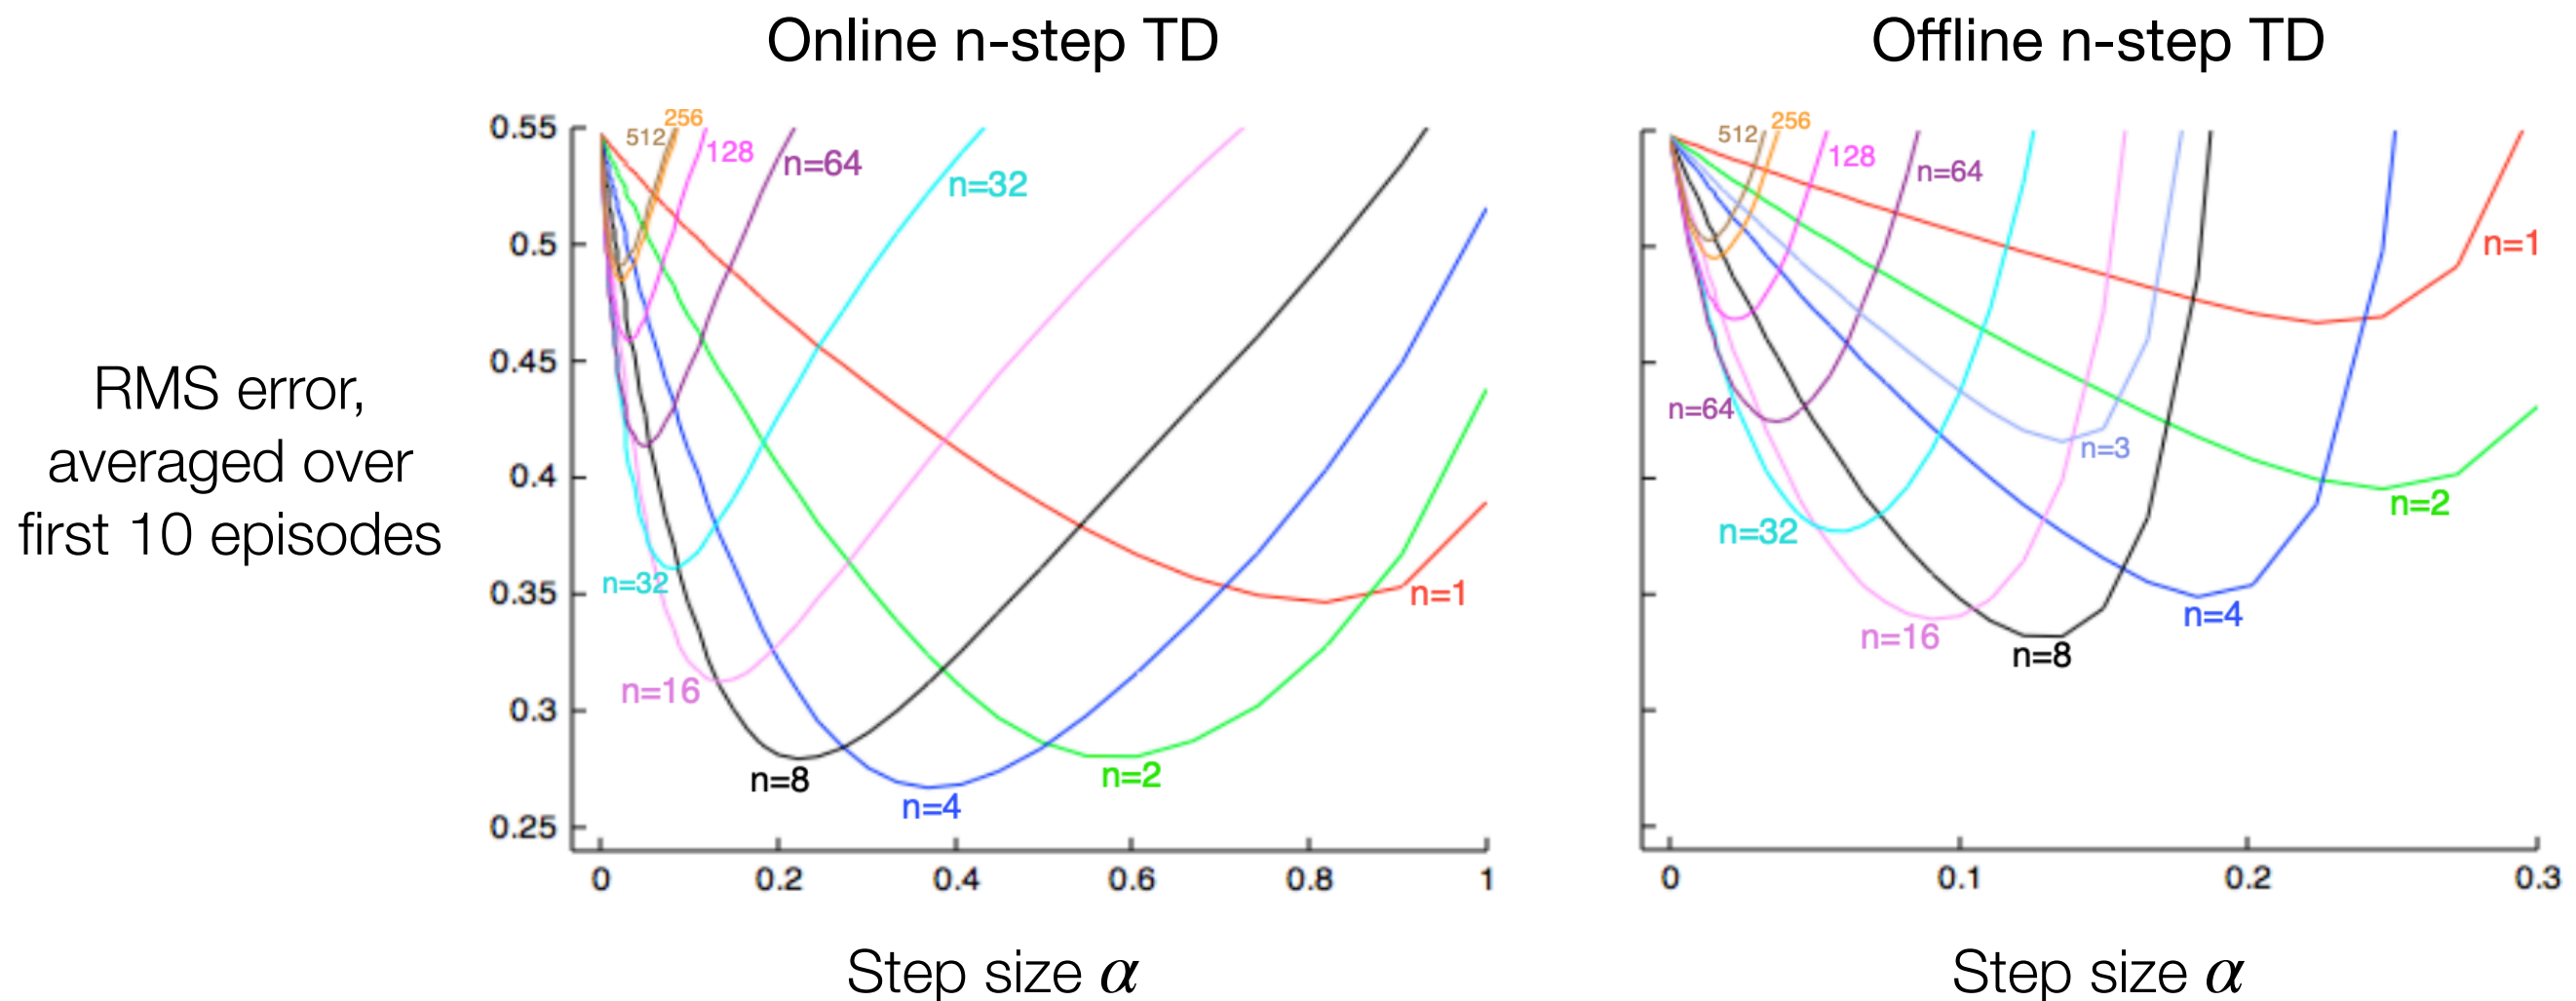


TD(0)	$n = 1$	$G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1})$
	$n = 2$	$G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+2})$
		\vdots
MC	$n = \infty$	$G_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$

Define the n-step return $G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$

n-step temporal-difference learning $V(S_t) \leftarrow V(S_t) + \alpha (G_t^{(n)} - V(S_t))$

Example - Large Random Walk



Observations

- Online methods generally worked best on this task, reaching lower levels of absolute error
- Methods with an intermediate value of n worked best
 - Generalization of TD and Monte Carlo methods to **n-step methods** can potentially perform better than either of the two extreme methods

Averaging n-Step Returns

Motivation

- Methods with an intermediate value of n worked best
 - How to pick the best n ?


Averaging n-step returns

- We can average n -step returns over different n
 - e.g., average the 2-step and 4-step returns

$$\frac{1}{2}G^{(2)} + \frac{1}{2}G^{(4)}$$

- Combines information from two different time-steps

Can we efficiently combine information from **all** time-steps?




λ -return


The λ -return G_t^λ combines all n-step returns $G_t^{(n)}$

- Using weight $(1 - \lambda)\lambda^{n-1}$: $G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$
- TD(λ) backups: $V(S_t) \leftarrow V(S_t) + \alpha(G_t^\lambda - V(S_t))$

TD(λ), λ -return


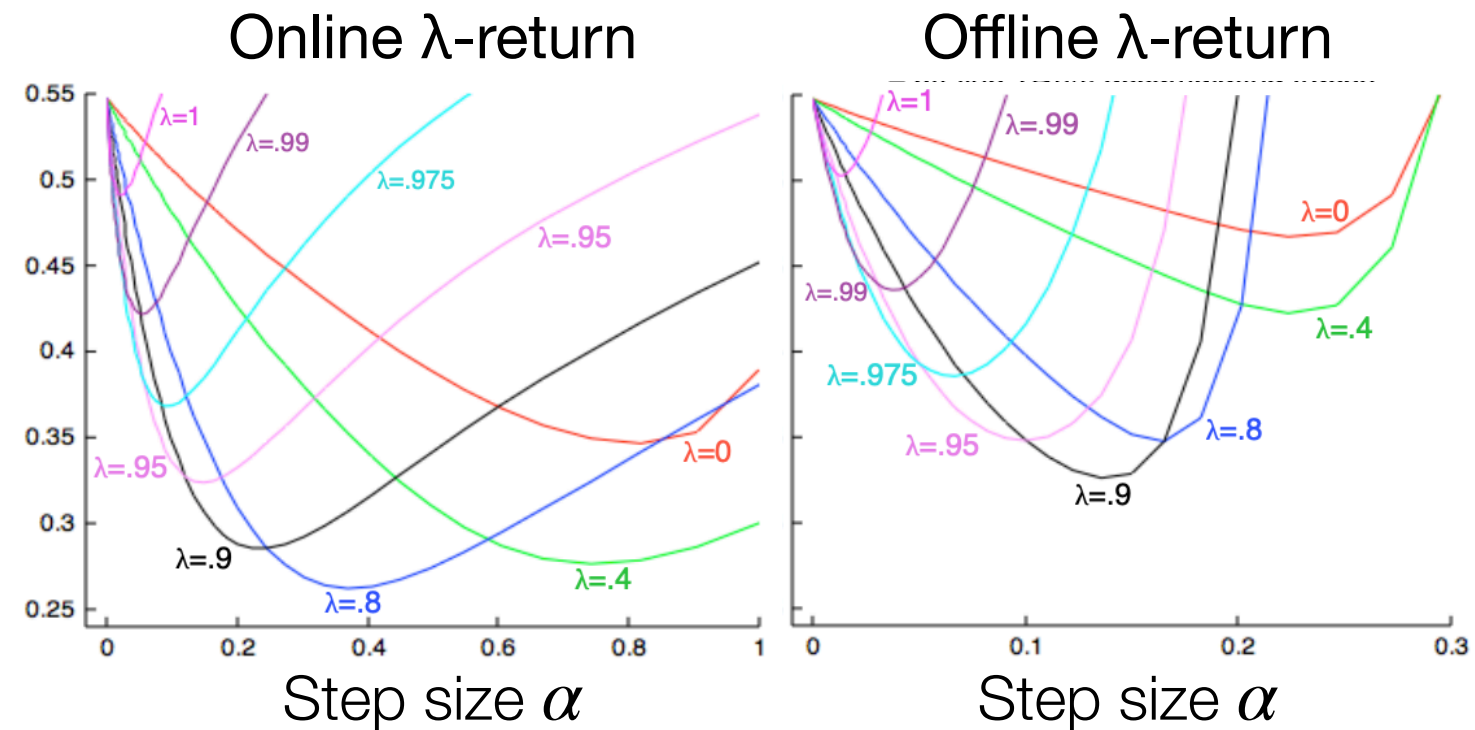


Weighting given in the λ -return to each of the n-step returns



Forward-view TD(λ)

RMS error,
averaged over
first 10 episodes



- Update value function towards the λ -return
- Forward-view looks into the future to compute G_t^λ
- Like MC, can only be computed from complete episodes

Backward-view TD(λ) and Eligibility Traces

Goal: update online, every step, from incomplete sequences


Eligibility Traces

- Credit assignment problem: did bell or light cause shock?
 - Frequency heuristic: assign credit to most frequent states
 - Recency heuristic: assign credit to most recent states
- Eligibility traces combine both heuristics



$$E_0(s) = 0$$

$$E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$

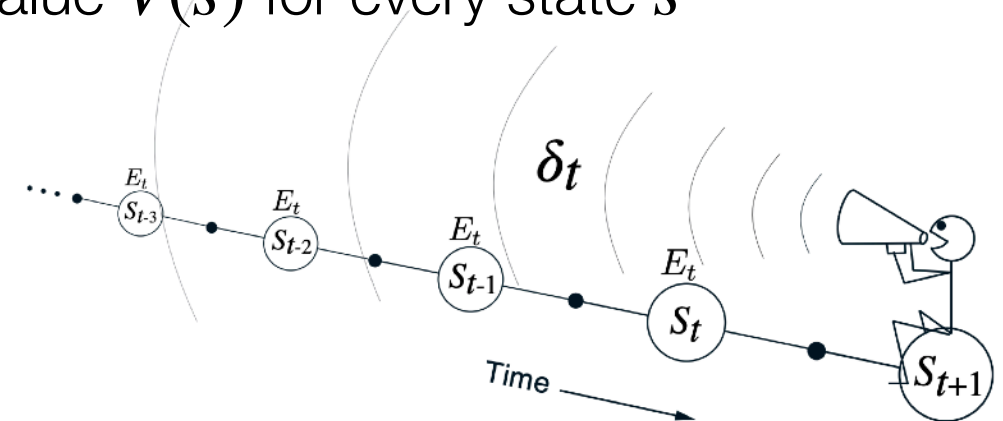


Backward-view TD(λ)

- Keep an **eligibility trace** for every state s and update value $V(s)$ for every state s
- In proportion to TD-error δ_t and eligibility trace $E_t(s)$

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$



TD(λ) and TD(0)

Theorem

- The sum of offline updates is identical for forward-view and backward-view TD(λ)

$$\sum_{t=1}^T \alpha \delta_t E_t(s) = \sum_{t=1}^T \alpha (G_t^\lambda - V(S_t)) \mathbf{1}(S_t = s)$$

$\lambda=0$

- Only current state is updated $E_t(s) = \mathbf{1}(S_t = s)$ and $V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$
- This is exactly equivalent to TD(0) update $V(s) \leftarrow V(s) + \alpha \delta_t$

MC and TD(1)

$\lambda=1$

- Credit is deferred until end of episode
- Consider episodic environments with offline updates
- Over the course of an episode, total update for TD(1) is the same as total update for MC

Consider an episode where s is visited once at time-step k ,

- TD(1) eligibility trace discounts time since visit,

$$E_t(s) = \gamma E_{t-1}(s) + \mathbf{1}(S_t = s) = \begin{cases} 0, & \text{if } t < k \\ \gamma^{t-k}, & \text{if } t \geq k \end{cases}$$

TD(1) updates accumulate error online

$$\sum_{t=1}^{T-1} \alpha \delta_t E_t(s) = \alpha \sum_{t=k}^{T-1} \gamma^{t-k} \delta_t = \alpha (G_k - V(S_k))$$

By the end of episode it accumulates total error

$$\delta_k + \gamma \delta_{k+1} + \gamma^2 \delta_{k+2} + \dots + \gamma^{T-1-k} \delta_{T-1}$$

MC and TD(1)

When $\lambda = 1$, sum of TD errors telescopes into MC error,

$$\begin{aligned} & \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \dots + \gamma^{T-1-t}\delta_{T-1} \\ &= R_{t+1} + \gamma V(S_{t+1}) - V(S_t) \\ & \quad + \gamma R_{t+2} + \gamma^2 V(S_{t+2}) - \gamma V(S_{t+1}) \\ & \quad + \gamma^2 R_{t+3} + \gamma^3 V(S_{t+3}) - \gamma^2 V(S_{t+2}) \\ & \quad \dots \\ & \quad + \gamma^{T-1-t} R_T + \gamma^{T-t} V(S_T) - \gamma^{T-1-t} V(S_{T-1}) \\ &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-1-t} R_T - V(S_t) \\ &= G_t - V(S_t) \end{aligned}$$

$\lambda=1$

- TD(1) is roughly equivalent to every-visit Monte-Carlo
 - Except that error is accumulated online, step-by-step
- If value function is only updated offline at end of episode
 - Then total update is exactly the same as MC

Equivalence of Forward and Backward TD(λ)

Consider an episode where s is visited once at time-step k ,

- TD(λ) eligibility trace discounts time since visit,

$$E_t(s) = \gamma E_{t-1}(s) + \mathbf{1}(S_t = s) = \begin{cases} 0, & \text{if } t < k \\ (\gamma\lambda)^{t-k}, & \text{if } t \geq k \end{cases}$$

Backward TD(λ) updates accumulate error online

$$\sum_{t=1}^T \alpha \delta_t E_t(s) = \alpha \sum_{t=k}^T (\gamma\lambda)^{t-k} \delta_t = \alpha (G_k^\lambda - V(S_k))$$

By end of episode it accumulates total error for λ -return

For multiple visits to s , $E(s)$ accumulates many errors



Equivalence of Forward and Backward TD(λ)

Offline updates

- Updates are accumulated within episode
 - but applied in batch at the end of episode

Online updates

- TD(λ) updates are applied online at each step within episode
- Forward and backward-view TD(λ) are slightly different
- Seijen, Harm, and Rich Sutton. "[True online TD \(lambda\)](#)." *ICML*, 2014.
 - Exact online TD(λ) achieves perfect equivalence
 - By using a slightly different form of eligibility trace



Summary of Forward and Backward TD(λ)

Offline updates	$\lambda=0$	$\lambda \in [0, 1]$	$\lambda=1$
Backward view	TD(0)	TD(λ)	TD(1)
Forward view	TD(0)	Forward TD(λ)	MC
Online updates	$\lambda=0$	$\lambda \in [0, 1]$	$\lambda=1$
Backward view	TD(0)	TD(λ)	TD(1)
Forward view	TD(0)	Forward TD(λ)	MC
Exact Online	TD(0)	Exact Online TD(λ)	Exact Online TD(1)


= here indicates equivalence in total update at end of episode

Summary - Model-Free Prediction

Monte-Carlo backup

- Based on the **entire sequence** of observed rewards until the end of the episode


$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$



Temporal-Difference Backup

- Based on just the one next reward, using the value of the **state one step later** as a proxy for the remaining rewards

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



TD(λ) Backup

- Based on an **intermediate number of rewards**: more than one, but less than all of them until termination

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t^\lambda - V(S_t))$$

